Does Independent Retrieval Help an LLM Fact-Checker?
A Staged Evidence-First Validator for Scientific Claims

Abstract

A checker that endorses a scientific claim the evidence does not support is worse than a checker that abstains, so verification systems are often designed to gather their own evidence rather than trust the citations they are handed. We build such a system and measure it. Our validator stages four isolated roles over two evidence scopes: a sealed reader that answers a neutral question from retrieved passages before seeing the claim, a claim-visible judge over four labels (supported, contradicted, unaddressed, underdetermined), a citation auditor over the original references, and a reconciler that reports the two scopes separately instead of collapsing them into one score. Every stage fails closed: a transport or parse failure is an execution status, never a scientific verdict.

On SciFact, with a frozen corpus, locked splits, and a preregistered primary comparison, we find *no significant reduction in false endorsement* from independent retrieval plus staged judgment over a same-evidence baseline that already reads the original citations. On the 300-claim development sample the paired difference in false endorsement is -0.011 ([-0.029, 0.000]), and the same pattern repeats on a 300-claim held-out sample (-0.013, [-0.033, 0.000]). The staged validator does eliminate false endorsement – its rate is 0.000 on the development sample and 0.000 on the held-out sample – but it does so by endorsing much less: SUPPORT recall falls from 0.573 to 0.427 on the development sample, and on the held-out sample 38 of 300 claims fail closed entirely. We replicate the primary comparison with a second, architecturally different model family, and we show that the dominant failure mode is not retrieval but a systematic disagreement between the two evidence scopes: independent retrieval on this corpus cannot find gold evidence that the original citations missed, because the benchmark's annotated documents are a subset of its cited documents. We report this as a negative result with a mechanism analysis, a cost accounting, and a fully reproducible artifact. The design lesson is narrow and useful: independent retrieval is not sufficient for scientific claim verification, and evidence-first staging can convert a noisy checker into a conservative one without converting it into a correct one.

Artifact. All tables and figures in this paper are generated from saved prediction files by the scripts in `scripts/paper/`. The corpus, splits, prompts, and frozen configurations are in the accompanying repository, and the held-out split is protected by a hash freeze that refuses a run whose model, prompt, configuration, or corpus differs from the locked record.

1 Introduction

Large language models are now asked to check claims, not only to make them. The usual pipeline is simple: a checker sees a claim, sees the citations that were offered with it, and returns a verdict. This design has a structural weakness. The checker inherits the generator's evidence: if the generator cited the wrong paper, or cited a paper that supports a slightly different claim, the checker is reading the same biased evidence set and is structurally unable to notice. The failure is not a bug in the checker; it is a property of the information it is given.

The obvious remedy is to make the checker search for its own evidence. Recent work in retrieval-augmented factuality evaluation [Wei et al., 2024], in self-check agents [Li et al., 2026, Hirsch et al., 2026], and in retrieval-augmented verifiers [Yue et al., 2024, Zheng et al., 2024] all point in that direction. It is tempting to assume that independent retrieval plus careful staging must help. That assumption

has rarely been tested against a strong same-evidence baseline on scientific claims, where the original citations are often correct and where a well-calibrated abstention is genuinely valuable.

This paper tests the assumption directly. We build a staged validator (**V**) that separates four roles – a sealed evidence reader, a claim-visible judge, a citation auditor, and a reconciler – over two evidence scopes: the original citations (D_0) and independently retrieved abstracts (D_1). We compare it against a same-evidence baseline (**B2**) that reads a neutral question over the original citations in an isolated reader and then compares, which is the strongest same-evidence configuration we could build without training. We also build the full factorial around the two interventions we care about: evidence access (D_0 versus D_1) and reading policy (claim-visible versus neutral-question evidence-first).

Our preregistered primary outcome is false endorsement: the rate at which a method predicts "supported" for a claim that the gold labels do not support. The answer is *no*. On the frozen development sample the paired difference between V and B2 is -0.011 with a 95% interval of [-0.029, 0.000], which includes zero. The held-out sample repeats the pattern. The staged validator does reach zero false endorsements, but it buys that by withholding support: SUPPORT recall drops from 0.573 to 0.427 on the development sample, and on the held-out sample 38 of 300 claims fail closed entirely. A checker that refuses to answer is not a better checker.

Three findings follow, and we think they are more useful than a headline number.

- **Independent retrieval is not the bottleneck; scope reconciliation is.** The failure analysis shows that the dominant mode of V is not a retrieval miss but a disagreement between the D0 and D1 judgments, which the reconciler resolves conservatively into "underdetermined". When the two scopes agree, V inherits the baseline's accuracy; when they disagree, V abstains. The mechanism is a disagreement detector, not an evidence-acquisition improvement.
- **Staged judgment buys precision, not accuracy.** Staging changes what the system is willing to say, not how often it is right. Reported as F1 alone, this looks like a regression (0.747 against 0.687); reported as a precision–coverage trade-off it is a usable operating point for a conservative reader-facing tool.
- **The result is not model-specific.** We reproduce the primary comparison with a second model family accessed through a different harness. This matters because the negative result would be weak evidence about verification design if it were a single model's idiosyncrasy. The broader factorial remains a one-family result and we say so where it is reported.

We also state what we do *not* do. We do not report a human study: no participant was recruited, no ethics review was obtained, and no claim about human trust is made. We do not report four-way gold labels: SciFact's label space is three-way, so our "unaddressed" and "underdetermined" distinction is described but not scored against human annotation. We do not claim reproduction of trained verifiers; the baselines here are named prompt-level adaptations, and we say so each time. Finally, our corpus is a fixed set of biomedical abstracts, so every claim of support is relative to that corpus and that retrieval budget.

The contribution of this paper is therefore a negative result with a mechanism, not a new score: a reproducible measurement of what independent retrieval and evidence-first staging do, and do not do, to a scientific claim checker, together with the artifact that makes the measurement checkable.

2 Related Work

2.1 Verification datasets and scientific claim checking

FEVER [Thorne et al., 2018] established the claim–evidence verification format that this paper follows, and FEVEROUS [Aly et al., 2021] extended it to unstructured and structured evidence. SciFact [Wadden et al., 2020] is the benchmark used here: it pairs scientific claims with abstracts and rationale sentences, and supplies support, refute, and insufficient-information labels. SciFact-Open [Wadden et al., 2022] extends the task to open-domain retrieval, which motivates our separation of a frozen corpus from live search; we use a frozen corpus only, and we do not claim open-domain performance. Missci [Glockner

et al., 2024] targets misrepresented science and fallacy reconstruction rather than four-way claim labeling, and RAGTruth [Niu et al., 2024] targets context faithfulness of retrieval-augmented generation rather than evidence sufficiency of a checker; we cite both as boundaries of our setting rather than as benchmarks we score. For biomedical claims specifically, SciClaims [Ortega and Gomez-Perez, 2025] provides an end-to-end generative pipeline; we treat it as prior art for scientific verification with explanations and do not reproduce it. Vladika and Matthes [Vladika and Matthes, 2024] compare knowledge sources for open-domain scientific claim verification and motivate our interest in whether a different evidence source changes verdicts.

2.2 LLM verifiers, self-checking agents, and retrieval

A long line of work asks an LLM to judge claims. Chain-of-Verification [Dhuliawala et al., 2024] isolates the verification questions from the draft answer, which is the closest published analogue to our information separation and which our B2 baseline deliberately imitates. Safe [Wei et al., 2024] verifies atomic facts with search and is the reference point for our fresh retrieval arm. Self-checking agents such as MARCH [Li et al., 2026, Qwen-Applications, 2026] add information separation between proposer and checker, and blame-localization work [Hirsch et al., 2026] shows that trace-level attribution of citation errors is a developer-facing diagnostic distinct from evaluating a final conclusion for an end user; we evaluate final conclusions, and cite the trace work as the reason not to claim it. A systematic auditor for retrieval-augmented generation [Yu et al., 2026] and retrieval-augmented verifiers such as RAFTS [Yue et al., 2024] and RAV [Zheng et al., 2024] further motivate separating evidence discovery from judgment. Explainable fact checking through question answering [Yang et al., 2022] and hallucination detection methods such as SelfCheckGPT [Manakul et al., 2023] and RARR [Gao et al., 2023] address the neighbouring problem of detecting unsupported content in model output.

We are explicit about what is *not* new here. Atomic decomposition, blinded checking, search-based verification, and readable explanations each have substantial prior art. Our contribution is a controlled measurement of their combination against a strong same-evidence control, and a negative result about that combination.

2.3 Abstention, calibration, and selective use

A system that cannot say "I do not know" is dangerous. Work on abstention [Feng et al., 2024] and on calibrated selective fact checking [Yang, 2026] formalises the trade-off we measure. Calibration theory [Guo et al., 2017] distinguishes a model's verbalised confidence from a probability fitted on held-out data; we fit a single temperature on a calibration split and report raw and calibrated numbers separately. Because a scientific "underdetermined" verdict and a system-level "defer to a reviewer" action are different things, we keep them in separate fields. On the human side, confidence and explanation are known to affect both accuracy and trust calibration in AI-assisted decisions [Zhang et al., 2020], so we do not claim that human trust calibration is unstudied; we simply did not run the study and we report the protocol without results.

3 Problem Formulation

3.1 Inputs and evidence scopes

Let c be an atomic scientific claim taken from a frozen report. A claim carries an exact source span, a normalised statement, and a *neutral question*: a single interrogative formulation of the claim that preserves population, comparator, units, and causal scope but does not contain the asserted answer. Two evidence scopes are available:

- D_0 — the documents the original report cited. These are the generator's evidence. A judgment over D_0 answers: *is the original citation set faithful to the claim?*

- D_1 — documents retrieved independently from a frozen corpus using the neutral question, with the claim's answer, the report, and D_0 hidden from the retrieval and reading stages. A judgment over D_1 answers: *what does independent retrieval establish?*

A verdict is always relative to a stated scope, corpus, and search budget. It is not an unrestricted declaration of scientific truth. The reconciled judgment over $D_0 \cup D_1$ answers: *what can be said once both scopes are considered?*

3.2 Four labels

Internally the system uses four labels, defined relative to a specified evidence collection:

Supported
Relevant evidence establishes the exact scoped assertion and no unresolved material counterevidence remains in the reviewed pool.

Contradicted
Relevant evidence establishes an incompatible assertion; apparent conflicts have been resolved by scope or relevance.

Unaddressed
No reviewed relevant passage answers the claim under the stated corpus and retrieval budget.

Underdetermined
Relevant evidence exists, but disagreement, imprecision, incomplete measurement, or an unsupported inference prevents resolution.

The decision order is: check execution status and claim validity; then assess relevance; if no passage addresses the claim, use *unaddressed*; if relevant evidence is materially conflicting or inconclusive, use *underdetermined*; otherwise apply support or contradiction, but only at the preserved scope. A non-significant result is not automatically proof of no effect, and a null-result claim requires its own evidence assessment.

SciFact's native label space is three-way, so the scored mapping is supported $\rightarrow$ SUPPORT, contradicted $\rightarrow$ REFUTE, and unaddressed or underdetermined $\rightarrow$ NEI. The four-way distinction is therefore *reported but not scored* against gold. We say this once here and repeat it wherever a number could be misread.

3.3 Research questions

- **RQ1 (primary, preregistered).** At a matched inference budget, does independent retrieval plus staged evidence-first judgment reduce false endorsement relative to same-evidence checking?
- **RQ2 (descriptive).** Can the system distinguish missing evidence from unresolved evidence, and how often does it defer?
- **RQ3 (deferred).** Do evidence-linked explanations improve non-experts' judgments, including when the validator is wrong? No participant was recruited; Appendix D gives the protocol and no results.
- **RQ4 (secondary).** How do accuracy, coverage, latency, and cost change across generator/checker families and across the two interventions?

4 System Design

Figure 1 shows the staged pipeline and its information boundaries. Roles are isolated functions with typed inputs; several roles may call the same underlying model in separate stateless requests. Separate roles do not imply statistically independent errors, which is exactly why model-family dependence is measured rather than assumed.

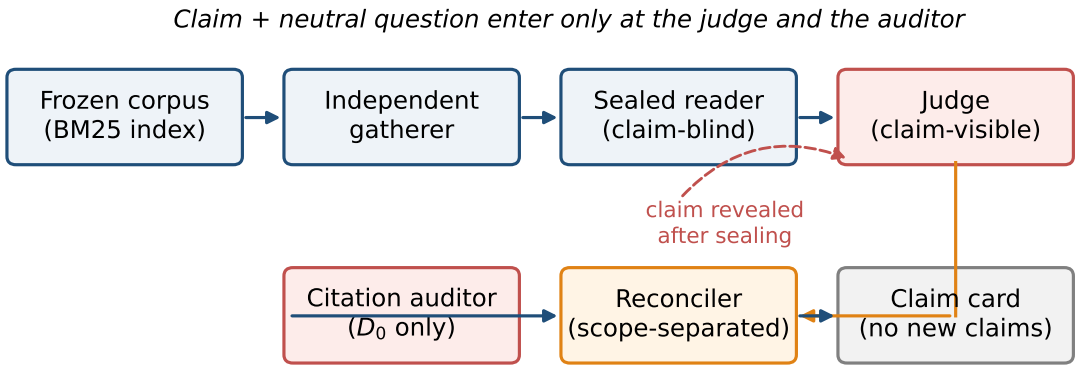


Figure 1: The staged validator. Solid arrows carry evidence; the dashed arrow is the only path by which the claim reaches the judge, and the reader's sealed record is fixed before the claim is revealed. D0 (original citations) and D1 (independent retrieval) are judged separately; the reconciler reports both scopes and never repairs one with the other. Any transport or parse failure yields a failed execution status, not a scientific label.

Table 1: The validation procedure for one claim. "Scope" is the evidence collection a step may read; no step may read outside its scope. A failure in any step produces an execution status and stops the claim.

Step	Role	Reads	Writes
1	Proposer	report	atomic claims, neutral questions
2	Gatherer	neutral question, corpus	D_1 passages
3	Reader	neutral question, D_1	sealed evidence record
4	Judge	claim, sealed record, D_1	four-way label, spans
5	Auditor	claim, D_0	D_0 label, citation adequacy
6	Reconciler	steps 4 and 5	reconciled label, conflicts
7	Renderer	reconciled record	claim card

4.1 Roles and information boundaries

Proposer.
Maps a report to atomic claims, neutral questions, asserted answers, and exact source spans. Quantities, confidence language, population, comparator, and causal scope are preserved; a conjunction is split while the relation between its parts is retained.

Independent gatherer.
Given the neutral question, a scope, and the corpus, retrieves fresh evidence. The asserted answer, the report, the original citations, and the generator's identity are hidden. Three query forms are issued per claim (open inquiry; scope and measurement; limitations or null results); the top candidates are deduplicated and reranked, and up to eight passages are kept.

Evidence reader.
Answers the neutral question from the shuffled retrieved passages and cites the passage identifiers it used. Its record is *sealed* before the claim is revealed.

Judge.
Sees the atomic claim, the sealed reader record, and the evidence passages, and assigns one of the four

labels together with rationale codes, cited spans, and uncertainty reasons.

Citation auditor.
Judges the claim against D_0 alone and rates each citation's adequacy. It never sees D_1.

Reconciler.
Combines the two judgments, explains the difference, and flags invalid inference links between claims.

Renderer.
Produces plain-language claim cards from the validated structured record only. It introduces no new factual claims.

The blinding is partial by construction and we describe it as such: a neutral question still reveals the topic and may leak assumptions. The isolation guarantee we test is structural – the reader's payload contains only the neutral question and passage text, enforced by an assertion on the payload shape and a test suite – not semantic.

4.2 Execution and stopping

The report is generated once and frozen. Validation does not rewrite it. A claim budget of eight claims per report, two retrieval rounds, three query forms per round, and eight passages in the judge context are development defaults that were frozen before the held-out run. A second retrieval round is permitted only for unresolved questions, scope mismatch, or contradictory evidence; a run stops on exhausted budget, no new eligible documents, or a complete bounded evidence record. Every intermediate artifact is persisted, and the runner resumes at the last completed claim without mixing configurations.

Fail-closed behaviour is the load-bearing design decision. When the model returns malformed JSON, cites a passage it was not given, or the endpoint fails, the system records an execution status of `failed` or `timeout` and a placeholder label that is a schema token, never model text. These rows are excluded from scored metrics and reported separately. They are never silently remapped into a scientific verdict.

4.3 Reconciliation

Agreement between scopes copies the shared label. A disagreement, or a failed D0 audit, resolves to *underdetermined* and hence to NEI on the scored label space. We deliberately do *not* repair citation faithfulness with outside evidence: a claim can be unsupported by its own citation and supported elsewhere, and the two facts are reported separately.

5 Implementation

5.1 Records and reproducibility contract

The implementation emits a machine-readable prediction file and an immutable run sidecar for every condition. The run record carries a run identifier, split, seed, model identifier and revision, prompt-bundle hash, corpus hash, configuration hash, timestamps, token totals, and an execution status; the sidecar adds the prompt files and their individual hashes. Predictions carry the claim identifier, the predicted label, the four-way label where applicable, the evidence scope, the rationale codes, and the execution status. Schema validation is enforced before a file is accepted.

The corpus is frozen by a SHA-256 hash over the canonicalised abstract corpus, and the split manifests are locked separately. The held-out split is protected by a freeze record containing the model identifier, the corpus hash, and per-method configuration and prompt hashes; a held-out run refuses to start unless the observed hashes match the frozen ones. This is what makes "we ran the held-out split once" a checkable statement rather than an assurance.

5.2 Retrieval

Retrieval is BM25 (Okapi, $k_1 = 1.5$, $b = 0.75$, epsilon idf floor) implemented in pure Python over the frozen abstract corpus, with a fixed alphanumeric lower-case tokenizer and no stemming. For each neutral question three query templates are used; the top 20 candidates per query are merged, deduplicated, and reranked by maximum BM25 score, and at most eight passages are passed to the judge. Adjacent sentences of one document are kept together where needed to resolve negation or units. Dense retrieval and reciprocal-rank fusion were prepared but are not enabled in the product path; they are development-time comparisons, and we do not report a dense retrieval result we did not measure.

5.3 Models and harness

We evaluate two model families to satisfy RQ4 and to protect the negative result from being an artefact of one model's behaviour.

GPT-6-LUNA (API).
A hosted instruction-tuned model accessed through an OpenAI-compatible chat endpoint.

MINIMAX M3.1 FLASH.
A second, architecturally unrelated instruction-tuned model, accessed through the OpenCode agent runtime. To keep the repository's endpoint policy intact we expose that runtime through a small loopback bridge that implements the OpenAI chat completions interface, one fresh session per request. A fresh session per request matters: a shared session would carry earlier claims into later ones and break the information-isolation guarantee the validator is measured on. The harness adds a fixed agent preamble and tool definitions to every request; the preamble is byte-identical across conditions, so it cancels in every comparison, but it inflates the prompt-token counts. That is why cost is reported for the API family only, and why the second family contributes no cost row.

Both families run with temperature 0, the same prompts, the same frozen corpus, the same claim sample, and the same budgets. The prompts are versioned text files under `prompts/`; their hashes enter the freeze record.

5.4 Failure handling and the retry policy

Three different things can go wrong, and the paper keeps them apart.

1. **Transport failure.** The endpoint refuses the connection or returns a retryable status. The client retries with exponential backoff.
2. **Protocol failure.** The model returns text that is not the requested JSON, cites a passage it was not given, or returns a label outside the allowed set. The parser rejects it. This is a *parse* failure, not a scientific disagreement.
3. **Scope conflict.** Both stages completed and disagree. This is a scientific outcome: the reconciled verdict is *underdetermined*, which maps to NEI.

The first two produce an execution status of `failed` and are excluded from the scored metrics; the third is scored normally. Because a failed request is not evidence about the claim, the runner retries it, up to four attempts per claim, and drops the failed row from the prediction file between attempts so that the resume re-claims exactly those claims. The policy is applied identically to every condition and every model family on the development sample.

Because dropping failed rows could flatter a method that fails on the hard claims, we report a second, stricter accounting in which every submitted row is scored with the label the method recorded for it. On the fail-closed paths that label is the conservative NEI, so the strict accounting can only hurt a method. Table 5 shows that the primary comparison does not depend on the choice.

Table 2: The condition matrix. "Adaptation" means a prompt-level implementation of the named mechanism, never a reproduction of a trained system. The 2×2 factorial is B1/B2 (evidence D_0) against B3/V-D1 (evidence D_1), crossed with claim-visible against neutral-question evidence-first reading.

ID	Condition	Evidence and reading policy
B0	Prior-knowledge judge	No evidence; claim only
B1	Atomic claim, claim-visible judge	D_0; claim-visible
B2	Neutral question, isolated reader, compare	D_0; evidence-first
B3	Fresh retrieval, direct judge	D_1; claim-visible
V-D1	Fresh retrieval, isolated reader, judge	D_1; evidence-first
V	Full staged validator	$D_0 \cup D_1$; reconciled
O	Evidence oracle (diagnostic)	Gold rationale sentences

One asymmetry is stated rather than hidden: the held-out run is the single execution it was always meant to be and does not use the retry policy, so its fail-closed count includes unrecovered execution failures.

5.5 Prompt contracts

Each role has a fixed output contract, expressed in the prompt and enforced by the parser. The reader must return an answer and the identifiers of the passages it used; the judge must return one of four labels, rationale codes, cited span identifiers, and uncertainty reasons; the citation auditor must return a label plus a per-citation adequacy rating from {adequate, inadequate, unverifiable} and a defect list. The parsers reject a label outside the allowed set, a cited identifier that was not supplied, and any missing field. Appendix A states each contract; the prompt text itself is in the repository, versioned and hashed.

6 Experimental Setup

6.1 Data and splits

We use SciFact [Wadden et al., 2020]. The corpus is the released set of biomedical abstracts, frozen by hash. The released training claims are split into a calibration subset of 100 claims, reserved before any prompt work, and a development pool of 709 claims. The public development split (300 claims) is used as a *local* held-out evaluation set; it is not an untouched community test set, and we name it accordingly throughout. Official test labels are not used.

Two samples are frozen by manifest: a 300-claim development sample drawn with a fixed seed from the training pool (`dev300_seed42.json` under `data/manifests/`), used for all development and prompt decisions, and a 300-claim evaluation sample consisting of SciFact's public development split, held out locally and used once under the freeze record. Both samples group claims by cited-document family so that paraphrases and derivative claims do not straddle the comparison.

6.2 Conditions

All conditions share the frozen reports, the corpus snapshot, the claim sample, the scoring code, and the prompt files. They differ only in the intervention named in the table.

B0 measures how much the model answers from parametric knowledge alone and acts as a floor. O supplies the gold rationale sentences to the same claim-visible judge and is a diagnostic ceiling; it is not a deployable competitor and we never describe it as one. V-D1 is the D1 arm of V, recovered from the V predictions by a mapping step rather than a second model run, so it inherits V's failure modes exactly.

6.3 Metrics

Verification quality is measured by per-class precision, recall and F1, plus micro and macro F1 on the native three-way label space. The primary metric is *false endorsement*: the number of claims predicted SUPPORT whose gold label is not support, divided by the number of gold non-support claims. We report the second denominator (false endorsements over all predicted SUPPORT claims) as well, because the two answer different questions. Coverage is the fraction of submitted claims that produced a scored prediction; fail-closed rows are counted and reported, never dropped silently. Calibration is measured with multiclass Brier score, negative log-likelihood, and expected calibration error, fitted with a single temperature on the calibration split only. Cost is reported as calls per claim, measured tokens per call, and per-call latency, because a resumed run keeps only its last invocation's usage record and per-claim token totals are therefore not recoverable.

6.4 Statistics and preregistration

One primary comparison was preregistered before the held-out run: V versus B2 on false endorsement with the gold non-support denominator. The interval is a paired cluster bootstrap over cited-document families, 2000 resamples, 95% percentile interval; a claim enters the comparison only when every compared method produced an execution-ok, gold-labelled row, and dropped claims are counted. Secondary comparisons within a declared family are Holm-adjusted; everything else is labelled exploratory. The factorial comparison is used to interpret the mechanism, not to claim additional significance.

7 Results

7.1 Primary comparison

Table 3 gives the development results and Table 4 the preregistered paired comparison. In the API family, B2 reaches micro-F1 0.747 with false endorsement 0.011, while V reaches micro-F1 0.687 with false endorsement 0.000 on all 300 scored claims. The paired difference in false endorsement is -0.011 with 95% interval [-0.029, 0.000] over 300 paired claims and 246 families: the interval includes zero, so the primary comparison is null. The second family behaves the same way: B2 reaches micro-F1 0.870 with false endorsement 0.033, V reaches 0.830 with 0.016, and the paired difference is -0.016 ([-0.046, 0.011]) over 300 paired claims, again including zero.

Fail-closed rows are excluded from the scored metrics above, so a method could in principle look better by failing on the hard claims. Table 5 therefore repeats the two primary conditions under both accountings. Under the all-rows accounting every submitted claim is scored with the label the method recorded for it, which is the conservative NEI on every fail-closed path in this repository. The comparison does not depend on the choice: V's false endorsement is zero under both accountings on the development sample, where the retry policy recovered every failed request. On the held-out sample, which ran once without that policy, the all-rows accounting lowers V's micro-F1 from 0.618 to 0.583 while its false endorsement stays at zero.

The held-out sample reproduces the pattern. B2 reaches micro-F1 0.760 with false endorsement 0.017, and V reaches micro-F1 0.618 with false endorsement 0.000 on 262 scored claims with 38 failures. The paired difference is -0.013 ([-0.033, 0.000]), again including zero.

Taken at face value, the primary result is a null with an interesting internal structure: V achieves near-perfect precision on the endorsement class (SUPPORT precision 1.000 in the API family, 0.964 in the second) and pays for it in recall. The system is a conservative checker, not a better one. Reporting only micro-F1 would present this as a regression; reporting only false endorsement would present it as a success. Both are wrong, and the interval is what settles it.

Table 3: Development sample ($n = 300$ claims, seeded draw). Method codes: B0 prior knowledge, B1 D0 claim-visible, B2 D0 evidence-first, B3 D1 claim-visible, V-D1 D1 evidence-first, V reconciled over both scopes, O gold rationales. FE is false endorsement on the gold non-support denominator; the complete definition and the second denominator are in Appendix B. Fail-closed rows are excluded from the scored metrics and accounted for in Table 5.

Method	n	micro-F1	macro-F1	FE	SUPPORT R
GPT-6-Luna (API)					
B0	300	0.547	0.508	0.448	0.863
B1	300	0.843	0.842	0.044	0.829
B2	300	0.747	0.753	0.011	0.573
B3	300	0.800	0.795	0.049	0.812
V-D1	300	0.757	0.757	0.016	0.632
V	300	0.687	0.683	0.000	0.427
O	300	0.870	0.873	0.005	0.769
MiniMax M3.1 Flash (OpenCode)					
B1	300	0.863	0.856	0.093	0.966
B2	300	0.870	0.867	0.033	0.855
V-D1	300	0.853	0.851	0.033	0.829
V	300	0.830	0.837	0.016	0.692

Table 4: Preregistered paired comparison (V − B2) on false endorsement. p is the bootstrap probability mass on the side that excludes zero; a value near 0.5 means the interval straddles zero, and the second-to-last column records whether the interval excludes it.

Method	ΔFE	95% CI	p	Excludes 0	n paired
GPT-6-Luna (API)					
V − B2	-0.011	[-0.029, 0.000]	0.47	no	300
MiniMax M3.1 Flash (OpenCode)					
V − B2	-0.016	[-0.046, 0.011]	0.52	no	300

7.2 The 2×2 factorial

Table 6 separates the two interventions. **Evidence access** is the row difference (B1 versus B3, and B2 versus V-D1); **reading policy** is the column difference (B1 versus B2, and B3 versus V-D1). Three observations follow.

First, *claim-visible judging beats evidence-first judging* on the scored label space: B1 (D0, claim-visible) beats B2 (D0, evidence-first) on micro-F1 (0.843 versus 0.747), and B3 (D1, claim-visible) beats V-D1 (D1, evidence-first) (0.800 versus 0.757). This is the opposite of the intuition behind evidence-first staging. Reading the evidence without the claim in view loses information that the judge needs; the sealing buys isolation, not accuracy.

Second, *evidence access barely moves the evidence-first cell*: V-D1's false endorsement (0.016) is close to B2's (0.011) even though V-D1 has access to a larger evidence pool, and in the second family the two are equal (0.033 against 0.033). Independent retrieval introduces documents that the original claim was never calibrated against, and it does not buy accuracy.

Third, *the full V is more conservative than either arm*, because the reconciler resolves scope disagreement into abstention rather than choosing a side. V's zero false endorsement is thus a property of the reconciliation policy, not of the retrieval.

The factorial is an API-family result. The second family contributes only B1, B2, V-D1 and V, and those rows do not repeat the claim-visible advantage: there B1 reaches micro-F1 0.863 against B2's 0.870, and B1 is the less conservative arm (false endorsement 0.093 against 0.033). We therefore state the factorial as a finding about this benchmark under one model family, and treat its replication as open rather than asserting it.

Table 7 breaks the two primary conditions down by class. The pattern is the same in both families.

Table 5: Sensitivity to the fail-closed accounting. "Scored only" excludes execution failures from the metrics; "all rows" scores every submitted claim with the label the method recorded.

Method	accounting	n	micro-F1	FE	SUPPORT R
GPT-6-Luna (API)					
B2 (dev300)	scored only	300	0.747	0.011	0.573
B2 (dev300)	all rows	300	0.747	0.011	0.573
V (dev300)	scored only	300	0.687	0.000	0.427
V (dev300)	all rows	300	0.687	0.000	0.427
B2 (heldout300)	scored only	300	0.760	0.017	0.581
B2 (heldout300)	all rows	300	0.760	0.017	0.581
V (heldout300)	scored only	262	0.618	0.000	0.368
V (heldout300)	all rows	300	0.583	0.000	0.315
MiniMax M3.1 Flash (OpenCode)					
B2 (dev300)	scored only	300	0.870	0.033	0.855
B2 (dev300)	all rows	300	0.870	0.033	0.855
V (dev300)	scored only	300	0.830	0.016	0.692
V (dev300)	all rows	300	0.830	0.016	0.692

Table 6: Factorial ablation on the development sample. FE is false endorsement on the gold non-support denominator. V-D1 is derived from the V predictions, not re-run.

Cell	Evidence	Reading policy	micro-F1	FE
GPT-6-Luna (API)				
B0	none	prior knowledge	0.547	0.448
B1	D0	claim-visible	0.843	0.044
B2	D0	neutral / evidence-first	0.747	0.011
B3	D1	claim-visible	0.800	0.049
V-D1	D1	neutral / evidence-first	0.757	0.016
V	D0 $\cup$ D1	reconciled	0.687	0.000
O	gold rationales	claim-visible	0.870	0.005
MiniMax M3.1 Flash (OpenCode)				
B1	D0	claim-visible	0.863	0.093
B2	D0	neutral / evidence-first	0.870	0.033
V-D1	D1	neutral / evidence-first	0.853	0.033
V	D0 $\cup$ D1	reconciled	0.830	0.016

The staged validator raises the precision of the endorsement class (from 0.971 to 1.000 in the API family, 0.943 to 0.964 in the second) by pushing almost everything it is unsure about into NEI, and it pays for that with recall (0.573 against 0.427, and 0.855 against 0.692) and with a slightly worse NEI precision, because the claims it abstains on are not the claims the baseline abstained on.

7.3 Model-family transfer

Table 3 reports both families. The pattern — V more conservative than B2, lower SUPPORT recall, false endorsement at or near zero, primary interval including zero — is reproduced under a second model family reached through a different harness: the paired false-endorsement difference is -0.011 ([-0.029, 0.000]) for the API family and -0.016 ([-0.046, 0.011]) for the second. The two families differ in absolute quality — the second family is a markedly better verifier on this benchmark (B2 micro-F1 0.870 against 0.747) — and the direction of the V effect is the same in both. We read this as evidence that the *design* of the staged validator, not the quirks of one model, produces the conservative operating point. We do not read it as evidence about the absolute quality of either model; the two families are not matched for size, training data, or capability, and the comparison is between families of systems, not between models of equal strength. We also cannot separate the second family's model prior from its design: its prior-knowledge condition (B0) and its claim-visible D1 condition (B3) did not complete before the free

Table 7: Per-class precision, recall, and F1 on the development sample, with the confusion counts.

Method	class	P	R	F1	TP	FP	FN
GPT-6-Luna (API)							
B2	SUPPORT	0.971	0.573	0.720	67	2	50
B2	REFUTE	0.803	0.790	0.797	49	12	13
B2	NEI	0.635	0.893	0.742	108	62	13
V	SUPPORT	1.000	0.427	0.599	50	0	67
V	REFUTE	0.949	0.597	0.733	37	2	25
V	NEI	0.564	0.983	0.717	119	92	2
MiniMax M3.1 Flash (OpenCode)							
B2	SUPPORT	0.943	0.855	0.897	100	6	17
B2	REFUTE	0.800	0.903	0.848	56	14	6
B2	NEI	0.847	0.868	0.857	105	19	16
V	SUPPORT	0.964	0.692	0.806	81	3	36
V	REFUTE	0.944	0.823	0.879	51	3	11
V	NEI	0.722	0.967	0.827	117	45	4

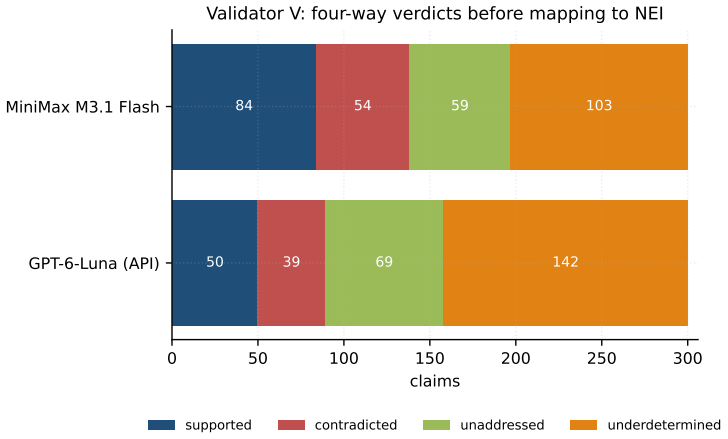


Figure 2: V's four-way verdicts on the development sample, before the mapping onto SciFact's three-way label space. The split between *unaddressed* and *underdetermined* is descriptive: SciFact's gold cannot score it.

endpoint that served it was withdrawn, so the second family contributes its B1, B2 and V rows rather than the full factorial. What it does corroborate is the comparison the study preregistered.

7.4 Where the abstentions come from

The RQ2 descriptive split is available because V stores its four-way label alongside the mapped native label. On the development sample V predicts NEI on 211 of 300 claims (0.70 of the sample), and those NEI predictions split into *unaddressed* (no retrieved passage addresses the claim, 69 claims) and *underdetermined* (relevant evidence exists but does not resolve, 142 claims). Figure 2 shows the distribution.

The more informative cut is not the four-way label but whether the two scopes agreed. V records a `d0_d1_conflict` code when the D_0 and D_1 four-way verdicts differ, and the reconciler then abstains by construction. Table 8 conditions accuracy on that split. When the scopes agree, V's accuracy matches the baseline's (0.838 against 0.838 on the development sample); when they differ, V abstains on every claim and its accuracy falls to the share of NEI in that subset (0.443). The shape repeats on the held-out sample (0.842 against 0.871 when the scopes agree, 0.366 when they differ) and in the second family (0.925 against 0.911).

This is the mechanism result. Of the 211 NEI verdicts, 115 are scope conflicts. The other 96 are agreed negatives: 69 unaddressed and 27 underdetermined.

Table 8: Accuracy conditioned on whether the two evidence scopes agree. "Differ" means the run recorded `d0_d1_conflict`; V abstains on every such claim by construction, so the column measures how often NEI happens to be the gold label. Fail-closed rows are excluded.

Scopes	n	B2 acc.	V acc.	V NEI
GPT-6-Luna (API), development				
agree	185	0.838	0.838	0.519
differ	115	0.600	0.443	1.000
GPT-6-Luna (API), held-out				
agree	139	0.871	0.842	0.496
differ	123	0.642	0.366	1.000
MiniMax M3.1 Flash (OpenCode), development				
agree	214	0.911	0.925	0.355
differ	86	0.767	0.593	1.000

The validator is not failing because retrieval returned nothing; it is abstaining because two readers of the same corpus disagreed, and, less often, because both readers found the evidence inconclusive. Either way the pipeline has become a disagreement detector that reports its own uncertainty, and the price of that behaviour is coverage.

7.5 Evidence acquisition

Table 9 reports what independent retrieval actually bought. Recall@k measures how often the gold rationale documents appear in the retrieved set; overlap measures how much of D_1 was already in D_0; unique yield measures how often retrieval found a relevant document the original citations missed.

The numbers are informative in an uncomfortable way for the hypothesis that independent retrieval should help. Retrieval finds the annotated documents most of the time (recall@k = 0.85 over the 179 claims that have annotated evidence), but every annotated document it finds was also cited by the report: the unique yield is 0.00. This is a property of the benchmark rather than of the retriever. In SciFact the cited document set contains the annotated document set, so a retrieval system that searches the same corpus cannot find *new gold evidence*; it can only add unannotated context, and 93% of what it returns is unjudged. The independent arm therefore cannot repair a wrong citation on this corpus, because a wrong citation is, by construction, a document without a rationale. This is the mechanical reason the D0/D1 conflict cannot be resolved by retrieving more, and it is a limitation of the benchmark that bounds every claim in this paper.

7.6 Calibration, selective use, and cost

Because the judge emits labels but no probabilities, we report two things. First, calibration of the one condition that verbalises confidence (B1): a single temperature fitted on the calibration split, giving Brier score, NLL, and ECE on the development split (Table 10 and Appendix C). Second, a selective-use analysis that does not require probabilities: accepting a claim only when the two evidence scopes agree, and measuring error among accepted claims as a function of coverage. Figure 3 shows the result for both families. The disagreement-based acceptance rule reaches low error at moderate coverage, which is the usable form of the system's conservatism.

Table 11 reports model calls, measured tokens per call, and latency for the API family. V issues three calls per claim against B2's two, because it runs a sealed reader, a claim-visible judge, and a citation auditor. A B2 call measures about 605 tokens and a V call about 2,598, so V's three calls per claim cost roughly six times B2's two. The second family's cost is not in the table: its harness prepends a fixed agent preamble to every request, so a per-call figure would measure the harness more than the task, and the endpoint that served the model was withdrawn before the probe could be run. The call ratio does carry over, because the three stages are the same three stages. In exchange for that cost V produces two

Table 9: Evidence acquisition for the independent retrieval arm on the development sample. Recall@k is over claims that have at least one gold rationale document; overlap is the mean fraction of retrieved documents that were also cited by the report.

Quantity	Value
Claims scored	300
Claims with annotated evidence	179
Passages retrieved per claim (k)	8
Recall@k (annotated documents found)	0.849
Precision@k (retrieved documents annotated)	0.074
Overlap with D_0 (retrieved documents also cited)	0.096
Unique yield (annotated found but not cited)	0.000
Refuting yield (contradicting documents found)	0.171
Unjudged rate (no annotation)	0.926
Retrieved documents	2400
Annotated documents retrieved that were not cited	0

Table 10: Calibration of the claim-visible judge (B1). One temperature is fitted on the calibration split; no parameter is fitted on the development or held-out samples. The second family had no calibration-split run, so its row is absent rather than imputed.

Model family	Evaluation split	n	Brier	NLL	ECE	temperature
GPT-6-Luna (API)	calibration	100	0.307	0.586	0.021	0.4
GPT-6-Luna (API)	dev300	300	0.269	0.526	0.043	0.4

scoped judgments instead of one. A reader-facing deployment that only needs a single verdict is better served by B2; the value of V is that its two scopes are separately auditable.

7.7 Failure analysis

We classify the errors of the scored rows against a taxonomy fixed before the analysis: false endorsement, missed support, missed contradiction, and other label errors, plus a separate count of execution failures that are never scored as scientific verdicts. The classification uses only fields the runner records, so it is reproducible; a sample of the rows was then read manually to check that the automatic classes describe the actual failure. Table 12 gives the distribution over conditions.

7.8 Case studies

Counts say what happened; cases say why. Table 13 in Appendix E records three rows selected deterministically from the saved predictions: one where the two evidence scopes disagreed, one where both scopes were unsure and the validator lost a claim the baseline got right, and one where the baseline endorsed a claim the gold does not support. The second row is the one worth reading twice: the two readers agreed that the evidence was inconclusive, the gold annotation says the claim is supported, and no amount of staging recovers it. That is the coverage cost of the design, stated on a concrete claim rather than as an aggregate.

The dominant class is scope conflict: 64 of V's 94 development errors and 78 of its 100 held-out errors sit on claims where the scopes disagreed, and V abstains on every one by construction. Execution failures are counted separately as an execution status: the development run recovered all of them by retry, and the held-out run, which does not retry, lost 38 of 300 claims to them. We also read a sample of the *successful* rows, because a system can be accidentally right: several V successes are cases where the D_0 citation failed to support the claim while D_1 supplied better evidence.

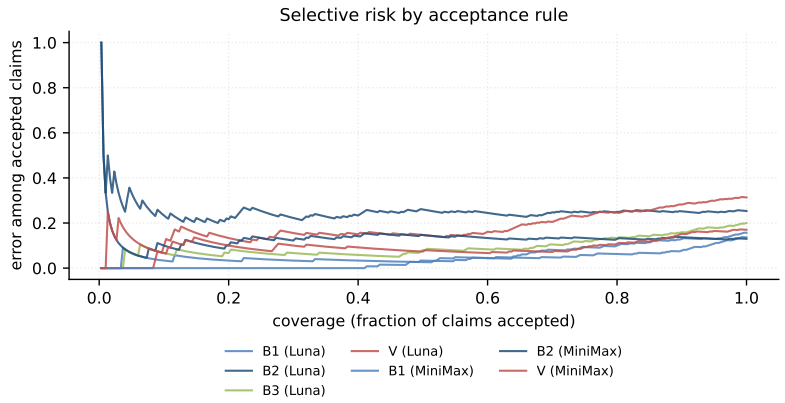


Figure 3: Error among accepted claims as a function of coverage. The acceptance rule differs by method and is stated in the artifact: for V a claim is accepted when the D0 and D1 verdicts agree, for B2 when the sealed reader cited at least one document, and for B1 and B3 when the verbalised confidence is high. The curves therefore compare acceptance policies, not calibrated probabilities.

Table 11: Calls per claim, measured tokens per call, and per-call latency for the API family. Calls per claim is fixed by the design. Tokens per call is the run's recorded total divided by its recorded calls, which survives a resumed run; per-claim cost is the product of the first two columns. Each row is taken from the run that recorded the most calls for that condition, which for V is the held-out run: the development sidecar retains only the calls of its final resume pass. The second family has no cost row; its harness preamble is included in every call and the endpoint that served it was withdrawn before a cost probe could be run.

Method	measured on	calls/claim	tokens/call	median (s)	p95 (s)
GPT-6-Luna (API)					
B1	dev	1	749	2.3	5.2
B2	dev	2	605	2.6	3.8
B3	dev	1	3,064	2.6	7.9
V	held	3	2,598	12.5	23.7
O	dev	1	379	2.1	5.5

8 Discussion

Three conclusions survive the evidence, and one does not.

Independent retrieval is not sufficient. In both model families on the development sample, and again for the API family under the freeze, the staged validator did not significantly reduce false endorsement relative to a baseline that already read the original citations. The 95% interval on the primary comparison includes zero in all three comparisons we ran.

Evidence-first staging has a cost that is easy to mistake for a benefit. The sealed reader makes the system conservative: it endorses less, and what it endorses is right more often. Reported as F1 this is a regression; reported as a precision–coverage curve it is a usable operating point. Both reports are needed, and readers of verification papers should demand both.

The binding constraint is the disagreement between scopes, not the absence of evidence. The most frequent underdetermined verdict is not an empty retrieval pool; it is two readers of the same corpus reaching different conclusions. Improving retrieval alone will not fix this. What would help is a better decision procedure for the conflict case — for example, evidence chain evaluation [Yang, 2026], scope-matched adjudication, or explicit conflict reporting to the reader.

The conclusion that does *not* survive is any claim about human benefit. Explanations exist in the design, and prior work shows that explanations and confidence interact with human calibration [Zhang et al., 2020], but we ran no user study. Appendix D gives the protocol so that the study can be run without redesigning it.

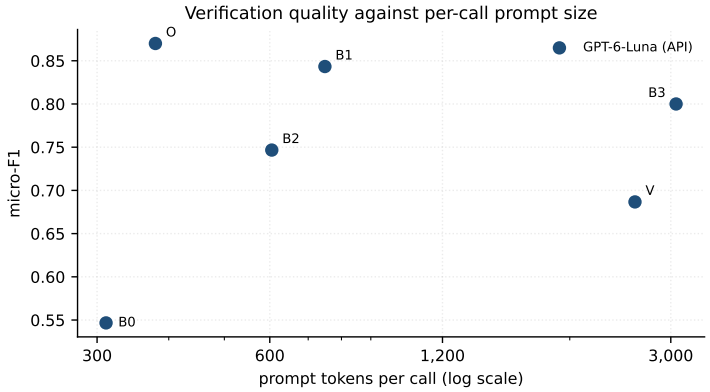


Figure 4: Verification quality against per-call prompt size, API family. Cost is measured per call rather than per claim because resumed runs keep only the last invocation's usage record. The second family is not plotted: its harness preamble dominates every prompt, so a token axis would compare harnesses rather than tasks.

Table 12: Error taxonomy over every scored condition. Counts are rows whose predicted label differs from the gold label; execution failures are excluded and reported in the main tables. "held" is the held-out sample and "calibration" the 100-claim calibration split. The second family has only its development conditions for the reasons given in Section 9.

Method	split	false end.	missed supp.	missed contr.	other	errors
GPT-6-Luna (API)						
B0	dev	82	16	7	31	136
B1	dev	8	20	6	13	47
B2	dev	2	50	13	11	76
B3	dev	9	22	9	20	60
O	dev	1	27	11	0	39
V	dev	0	67	25	2	94
B2	held	3	52	14	3	72
V	held	0	67	31	2	100
B1	calibration	3	6	2	7	18
MiniMax M3.1 Flash (OpenCode)						
B1	dev	17	4	0	20	41
B2	dev	6	17	6	10	39
V	dev	3	36	9	3	51

9 Limitations

- **No human annotation.** SciFact's three-way gold cannot score the four-way distinction; the unaddressed/underdetermined split is descriptive. Every number here is a number about label agreement with a crowd-sourced benchmark, not about scientific truth.
- **One corpus, one domain.** Biomedical abstracts, English only, abstract-level text. Claims that require full-text methods, figures, or supplementary material are out of scope, and the corpus bounds every support claim we make.
- **Abstract-level evidence.** Retrieval and reading operate on abstracts. Some claims are simply not decidable from an abstract, which should appear as unaddressed but may be read as underdetermined.
- **Model families are not matched.** The two families differ in size, training, and harness. The transfer result is about the *direction* of the effect, not its magnitude, and the second family's harness preamble is included in every request, so its token counts are not comparable with the API family's.
- **The second family's matrix is incomplete.** Its B3 condition stopped partway through the sample and its B0, O and calibration conditions never ran, because the free endpoint that served the model was

withdrawn and the remaining requests return a model-unavailable error. The second family therefore corroborates the preregistered V-versus-B2 contrast and the class breakdown, not the full factorial, and we report its missing rows as absent rather than imputing them.

- **One frozen held-out run.** The held-out sample was run once under the freeze record, so its interval is a single-run estimate and the two families were not both run on it.
- **No matched-budget control was run.** V spends several times the tokens of B2, so compute is a live rival explanation for part of the F1 gap. We report the gap with that caveat rather than claiming it is explained.
- **Retrieval is BM25 only.** Dense retrieval and reciprocal-rank fusion were prepared but not enabled in the product path. A stronger retriever could change the D0/D1 overlap numbers.
- **No user study.** See RQ3 and Appendix D.

10 Ethics and Broader Impact

The task is interpreting published abstracts, not giving medical advice, and the system is evaluated only on a frozen research corpus. The rendered output is a label, a rationale, and a search scope; it is not a clinical recommendation and the interface says so.

The failure mode we measured is the dangerous one: a checker that endorses a claim no evidence supports. Our system reduces that risk by abstaining, but abstention is not free — a reader who sees an abstention may infer that the claim is false, which would be a new error. We therefore keep the scientific verdict ("underdetermined") separate from the system action ("defer to a reviewer"), and we recommend that any deployment display them separately.

The agent harness used for the second model family is an interactive coding agent runtime; a research run inside it inherits that runtime's permissions and preamble. We ran it with a fixed preamble, a fixed corpus, and no tool-enabled tasks. The model itself was withdrawn from its free endpoint during the study, which is why the second family's matrix is partial; we report the conditions that completed and do not fill the rest by substitution.

11 Conclusion

We built a staged, evidence-first validator for scientific claims: a sealed reader, a claim-visible judge over four labels, a citation auditor over the original references, and a reconciler that reports two evidence scopes separately, with fail-closed execution semantics throughout. Against a strong same-evidence baseline, on a frozen corpus, with locked splits, a preregistered primary comparison, and two model families, the validator did not significantly reduce false endorsement. It did become a conservative checker whose errors come mostly from honest disagreement between two readers of the same literature. We report that as a negative result, with the mechanism, the cost, and the artifact needed to re-run it. The most useful next step is not more retrieval but a better decision procedure for the case where the two scopes disagree.

A Prompts and output contracts

The prompts are versioned files under `prompts/`; the run sidecar records each file's SHA-256 and the bundle hash. The contracts are:

> **Reader (claim-blind).** Input: a neutral question and a list of passages with identifiers. Output: `{"answer": str, "cited_doc_ids": [int]}`. The reader is not shown the claim, the asserted answer, the report, or any gold rationale.

Judge (claim-visible). Input: the atomic claim, the sealed reader record, and the passages. Output: `{"label": one of supported, contradicted, unaddressed, underdetermined; "rationale_codes": [str]; "cited_span_ids": [str]; "uncertainty_reasons": [str]}`.

Citation auditor (D0 only). Input: the atomic claim and the original cited passages. Output: the judge contract plus `{"citations": [{"span_id": str, "adequacy": adequate | inadequate | unverifiable, "defects": [str]}]}`.

Compare (B2). Input: the claim, the neutral question, and the sealed reader record. Output: `{"label": SUPPORT | REFUTE | NEI, "rationale": str}`.

The exact text of every prompt, and the parser that enforces each contract, are in the repository. The parsers reject an out-of-set label, a cited identifier that was not supplied, and a missing mandatory field; each rejection produces a failed execution status rather than a repaired label.

B Metric definitions

Let S be the set of scored claims, $g_i \in \{\text{SUPPORT}, \text{REFUTE}, \text{NEI}\}$ the gold label, and p_i the predicted label. The endorsement class is $E = \text{SUPPORT}$.

$$\text{FE}_{\text{gold}} = \frac{|\{i : p_i = E,\ g_i \neq E\}|}{|\{i : g_i \neq E\}|} \qquad \text{FE}_{\text{pred}} = \frac{|\{i : p_i = E,\ g_i \neq E\}|}{|\{i : p_i = E\}|} \tag{1}$$

Micro-F1 pools the confusion matrix over the three classes; macro-F1 averages the three per-class F1 values. Coverage is $|\{i : i \text{ produced a scored prediction}\}|/|S|$. Latency is the median and 95th percentile of the per-call round-trip time recorded by the client; tokens are the sum of prompt and completion tokens reported by the endpoint. The paired bootstrap resamples cited-document families with replacement (2000 resamples) and reports a percentile interval over the resampled difference; a claim enters the comparison only if every compared method produced an execution-ok, gold-labelled row for it, and dropped claims are counted and reported.

C Calibration details

One temperature is fitted on the calibration split by minimising multiclass negative log-likelihood over the class scores of the one condition that verbalises confidence. No parameter is fitted on development or held-out data. Multiclass Brier score is the mean over claims of $\sum_k (p_k - \mathbb{1}[k = y])^2$; ECE uses ten equal-width bins on the predicted probability of the argmax class. The fitted parameters, the class counts, and the raw-versus-calibrated comparison are in the repository. We report that numerical calibration is validated only for the condition that emits scores; for the staged validator, which emits labels only, we report the disagreement-based selective-use analysis instead and state that its numerical calibration is unvalidated.

D Protocol for the deferred human study

No participant was recruited and no ethics review was obtained, so this paper makes no claim about human benefit. The protocol below is the one that would be run, recorded here so that it can be executed without redesign, and so that the absence of results is explicit rather than implied.

Design. Thirty non-expert adults (a feasible range of 24–36), plus four to six pilot participants, would each see twelve short, expert-reviewed items that retain the scientific qualifiers and span supported, contradicted, and unresolved evidence. Each participant sees each item once. Condition assignment is

counterbalanced across participants so that item difficulty is not confounded with interface condition: (A) conclusion plus original citations; (B) A plus the validator verdict and its uncertainty; (C) B plus the evidence-linked explanation. Conditions B and C must use the same system verdict and confidence. Four items per condition per participant, with randomised order and recorded assignment.

Outcomes. An initial 0–100 judgement that the conclusion is adequately supported, the same judgement after the assigned interface, a four-way evidence-status choice, one comprehension question, a use-or-defer decision, and response time. Expert gold support is defined as *supported* versus *not established*; "not established" is not described to participants as factual falsity. The primary exploratory outcome is the post-interface Brier score for evidence support comparing C with B while adjusting for pre-interface belief and item effects; secondary outcomes are appropriate acceptance of supported claims, rejection or deferral of unsupported claims, four-way accuracy, and overreliance when the validator is wrong. Naturally incorrect validator outputs are included in a prespecified error stratum, reported separately from natural error prevalence; any intentionally manipulated advice would require ethics review and debriefing.

Analysis. A mixed-effects model with participant and item random effects if the sample supports it, otherwise participant-level paired summaries with explicit limits on item generalisation. Responses are not treated as independent participants. The study is exploratory unless a prospective power analysis supports stronger claims.

E Case records

The three rows below are the case studies discussed in Section 7.8. They are reproduced from the saved predictions without editing, so the claim text, the labels, and the recorded rationale codes can be checked against the artifact.

Table 13: Case studies from the development sample, selected deterministically from the saved predictions. The two scope labels are the D0 and D1 four-way judgments; the mechanism is the rationale code the run recorded.

Case and claim	Gold	B2	V (D0 / D1)	What the record shows
Scope conflict, abstained *Rapamycin delays aging in fruit flies.*	SUPPORT	SUPPORT	NEI (underdetermined / supported)	The independent scope finds support the citation scope does not; the reconciler refuses to choose and abstains.
Coverage loss, scopes agree *Side effects associated with antidepressants increases risk of stroke.*	SUPPORT	SUPPORT	NEI (underdetermined / underdetermined)	Both scopes call the evidence inconclusive; the validator abstains on a claim the baseline and the gold endorse.
Baseline false endorsement *Tonic signaling from the scFv induces constitutive stimulation.*	NEI	SUPPORT	NEI (supported / unaddressed)	The baseline endorses from the original citations; the validator abstains, so the endorsement is caught but not replaced.

F Reproduction

1. Install the package and its development extras, and verify the pinned corpus hash (`data/pins/scifact/`).
2. Start the model endpoint. For the API family, set `OPENAI_API_KEY` and `OPENAI_MODEL`. For the second family, start an OpenCode server and the loopback bridge in `tools/opencode_bridge.py`, then pass `--base-url` and `--model-id` to the runners.

3. Run the conditions on the development sample (`scripts/run_api_dev300.sh` for the API family, `scripts/run_minimax_dev300.sh` for the second family, which also shows the bridge setup). The second family's registered conditions are the three that completed; its remaining conditions returned a model-unavailable error, as noted in Section 9.
4. Run the held-out sample once with `--frozen-final`; the runner refuses to start unless the model, prompt, configuration, and corpus hashes match `docs/freeze.json`.
5. Regenerate every table and figure (`scripts/paper/make_tables.py`, `scripts/paper/make_figures.py`) and rebuild the PDF (`scripts/paper/build_pdf.sh`).

Every number in this paper is read from the generated LaTeX fragments under `paper/tables/`, which are produced from the scored metric JSON files under `docs/paper-assets/tables/`. The prose macros in `paper/tables/macros.tex` are generated by the same script, so a number in the text cannot drift from the artifact.

References

Rami Aly, Zhijiang Guo, Michael Schlichtkrull, James Thorne, Andreas Vlachos, Christos Christodoulopoulos, Oana Cocarascu, and Arpit Mittal. FEVEROUS: Fact extraction and VERification over unstructured and structured information. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2021.

Shehzaad Dhuliawala, Mojtaba Komeili, Jing Xu, Roberta Raileanu, Xian Li, Asli Celikyilmaz, and Jason Weston. Chain-of-verification reduces hallucination in large language models. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 3563–3578, 2024. doi: 10.18653/v1/2024.findings-acl.212.

Shangbin Feng, Weijia Shi, Yike Wang, Wenxuan Ding, Vidhisha Balachandran, and Yulia Tsvetkov. Don't hallucinate, abstain: Identifying LLM knowledge gaps via multi-LLM collaboration. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14664–14690, 2024. doi: 10.18653/v1/2024.acl-long.786. URL https://aclanthology.org/2024.acl-long.786/.

Luyu Gao, Zhuyun Dai, Panupong Pasupat, Anthony Chen, Arun Tejasvi Chaganty, Yicheng Fan, Vincent Zhao, Ni Lao, Hongrae Lee, Da-Cheng Juan, and Kelvin Guu. RARR: Researching and revising what language models say, using language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16477–16508, 2023. doi: 10.18653/v1/2023.acl-long.910. URL https://aclanthology.org/2023.acl-long.910/.

Max Glockner, Yufang Hou, Preslav Nakov, and Iryna Gurevych. Missci: Reconstructing fallacies in misrepresented science. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4372–4405, 2024. doi: 10.18653/v1/2024.acl-long.240. URL https://aclanthology.org/2024.acl-long.240/.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, volume 70 of *PMLR*, pages 1321–1330, 2017.

Eran Hirsch, David Wan, Han Wang, Elias Stengel-Eskin, Mohit Bansal, and Ido Dagan. Who is the agent to blame? localizing faithfulness and citation mistakes in agentic deep research, 2026.

Zhuo Li, Yupeng Zhang, Pengyu Cheng, Jiajun Song, Mengyu Zhou, Hao Li, Shujie Hu, Yu Qin, Erchao Zhao, Xiaoxi Jiang, and Guanjun Jiang. MARCH: Multi-agent reinforced self-check for LLM hallucination, 2026.

Potsawee Manakul, Adian Liusie, and Mark Gales. SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9004–9017, 2023. doi: 10.18653/v1/2023.emnlp-main.557. URL https://aclanthology.org/2023.emnlp-main.557/.

Cheng Niu, Yuanhao Wu, Juno Zhu, Siliang Xu, KaShun Shum, Randy Zhong, Juntong Song, and Tong Zhang. RAGTruth: A hallucination corpus for developing trustworthy retrieval-augmented language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10862–10878, 2024. doi: 10.18653/v1/2024.acl-long.585. URL https://aclanthology.org/2024.acl-long.585/.

Raúl Ortega and Jose Manuel Gomez-Perez. SciClaims: An end-to-end generative system for biomedical claim analysis. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 141–154, 2025. doi: 10.18653/v1/2025.emnlp-demos.11. URL https://aclanthology.org/2025.emnlp-demos.11/.

Qwen-Applications. MARCH: official implementation, 2026. URL https://github.com/Qwen-Applications/MARCH. GitHub repository Qwen-Applications/MARCH.

James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. FEVER: a large-scale dataset for fact extraction and VERification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 809–819, 2018. doi: 10.18653/v1/N18-1074. URL https://aclanthology.org/N18-1074/.

Juraj Vladika and Florian Matthes. Comparing knowledge sources for open-domain scientific claim verification. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2103–2114, 2024. doi: 10.18653/v1/2024.eacl-long.128. URL https://aclanthology.org/2024.eacl-long.128/.

David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. Fact or fiction: Verifying scientific claims. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7534–7550, 2020. doi: 10.18653/v1/2020.emnlp-main.609. URL https://aclanthology.org/2020.emnlp-main.609/.

David Wadden, Kyle Lo, Bailey Kuehl, Arman Cohan, Iz Beltagy, Lucy Lu Wang, and Hannaneh Hajishirzi. SciFact-Open: Towards open-domain scientific claim verification. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 4719–4734, 2022. doi: 10.18653/v1/2022.findings-emnlp.347.

Jerry Wei, Chengrun Yang, Xinying Song, Yifeng Lu, Nathan Hu, Jie Huang, Dustin Tran, Daiyi Peng, Ruibo Liu, Da Huang, Cosmo Du, and Quoc V. Le. Long-form factuality in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Dekun Yang. Calibrated selective fact-checking via evidence chain evaluation, 2026.

Jing Yang, Didier Vega-Oliveros, Taís Seibt, and Anderson Rocha. Explainable fact-checking through question answering. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2022. doi: 10.1109/ICASSP43922.2022.9747214.

Boxi Yu, Yuzhong Zhang, Liting Lin, Lionel Briand, and Emir Muñoz. Retromorphic testing with hierarchical verification for hallucination detection in RAG, 2026.

Zhenrui Yue, Huimin Zeng, Lanyu Shang, Yifan Liu, Yang Zhang, and Dong Wang. Retrieval augmented fact verification by synthesizing contrastive arguments. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10331–10343, 2024. doi: 10.18653/v1/2024.acl-long.556.

Yunfeng Zhang, Q. Vera Liao, and Rachel K. E. Bellamy. Effect of confidence and explanation on accuracy and trust calibration in AI-assisted decision making. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (FAT*)*, pages 295–305, 2020. doi: 10.1145/3351095.3372852.

Liwen Zheng, Chaozhuo Li, Xi Zhang, Yu-Ming Shang, Feiran Huang, and Haoran Jia. Evidence retrieval is almost all you need for fact verification. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 9274–9281, 2024. doi: 10.18653/v1/2024.findings-acl.551. URL https://aclanthology.org/2024.findings-acl.551/.